

Breast Cancer Classification

Distinguishing Malignant from Benign Tumors with Supervised Machine Learning

Course: CMP4336 – Introduction to Data Mining **Institution:** Bahçeşehir University, Faculty of Engineering and Natural Sciences

Group Members

Name	Student ID
Osman Yiğit Alver	2104312
İsacan Çeliktaş	2103746

Dataset: Breast Cancer Wisconsin (Diagnostic) — UCI ML Repository / Kaggle
(`uciml/breast-cancer-wisconsin-data`)

Table of Contents

1. Problem Definition
2. Data Exploration
3. Data Preprocessing
4. Model Development (Cross-Validation & Hyperparameter Tuning)
5. Model Evaluation
6. Interpretation of Results

Setup — Libraries

```
In [1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split, StratifiedKFold, cross_val_score
from sklearn.preprocessing import StandardScaler
from sklearn.pipeline import Pipeline
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from sklearn.metrics import (accuracy_score, precision_score, recall_score,
                             confusion_matrix, classification_report, roc_auc_score)

import warnings
warnings.filterwarnings('ignore')
```

```
RANDOM_STATE = 42
plt.style.use('seaborn-v0_8')
sns.set_palette('husl')
pd.set_option('display.max_columns', None)
print('Environment ready.')
```

Environment ready.

1. Problem Definition

Objective. Build and compare supervised machine-learning models that classify a breast tumor as **malignant (cancerous)** or **benign (non-cancerous)** from 30 quantitative features computed on digitized fine-needle-aspirate (FNA) images of breast masses.

Target variable. `diagnosis` — a binary label: **M (Malignant)** or **B (Benign)**. We encode it as **Malignant = 1** (positive class) and **Benign = 0**.

Why this matters. Breast cancer is a leading cause of cancer-related death worldwide, and early, accurate detection is decisive for survival. In a diagnostic setting a **false negative** (a malignant tumor predicted benign) is far more dangerous than a **false positive**. We therefore treat **recall on the malignant class** as the most important metric, alongside accuracy, F1-score and ROC-AUC.

Approach. We compare three supervised algorithms — **Logistic Regression** (linear baseline), **Random Forest** (bagged trees) and **Support Vector Machine** (RBF kernel) — using stratified 5-fold cross-validation, tune the two non-linear models with grid search, and evaluate the final models on a held-out test set.

2. Data Exploration

The dataset has **569 records** and **30 numerical features**. Each feature is reported as the **mean**, **standard error (se)** and **worst (largest)** value of ten cell-nucleus measurements (radius, texture, perimeter, area, smoothness, compactness, concavity, concave points, symmetry, fractal dimension).

```
In [2]: df = pd.read_csv('data.csv')
print('Raw shape:', df.shape)
df.head()
```

Raw shape: (569, 33)

Out [2]:

	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	sr
0	842302	M	17.99	10.38	122.80	1001.0	
1	842517	M	20.57	17.77	132.90	1326.0	
2	84300903	M	19.69	21.25	130.00	1203.0	
3	84348301	M	11.42	20.38	77.58	386.1	
4	84358402	M	20.29	14.34	135.10	1297.0	

In [3]: *# Column structure and data types*
df.info()

```
<class 'pandas.DataFrame'>
RangeIndex: 569 entries, 0 to 568
Data columns (total 33 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   id                                    569 non-null   int64
1   diagnosis                            569 non-null   str
2   radius_mean                          569 non-null   float64
3   texture_mean                         569 non-null   float64
4   perimeter_mean                       569 non-null   float64
5   area_mean                            569 non-null   float64
6   smoothness_mean                      569 non-null   float64
7   compactness_mean                     569 non-null   float64
8   concavity_mean                       569 non-null   float64
9   concave points_mean                  569 non-null   float64
10  symmetry_mean                        569 non-null   float64
11  fractal_dimension_mean               569 non-null   float64
12  radius_se                            569 non-null   float64
13  texture_se                           569 non-null   float64
14  perimeter_se                         569 non-null   float64
15  area_se                              569 non-null   float64
16  smoothness_se                        569 non-null   float64
17  compactness_se                       569 non-null   float64
18  concavity_se                         569 non-null   float64
19  concave points_se                    569 non-null   float64
20  symmetry_se                          569 non-null   float64
21  fractal_dimension_se                 569 non-null   float64
22  radius_worst                         569 non-null   float64
23  texture_worst                        569 non-null   float64
24  perimeter_worst                      569 non-null   float64
25  area_worst                           569 non-null   float64
26  smoothness_worst                     569 non-null   float64
27  compactness_worst                    569 non-null   float64
28  concavity_worst                      569 non-null   float64
29  concave points_worst                 569 non-null   float64
30  symmetry_worst                       569 non-null   float64
31  fractal_dimension_worst              569 non-null   float64
32  Unnamed: 32                          0 non-null     float64
dtypes: float64(31), int64(1), str(1)
memory usage: 146.8 KB
```

```
In [4]: # Target balance (raw labels)
print('Class counts:')
print(df['diagnosis'].value_counts())
print('\nClass proportions:')
print((df['diagnosis'].value_counts(normalize=True) * 100).round(1).astype(s
```

```
Class counts:
diagnosis
B    357
M    212
Name: count, dtype: int64
```

```
Class proportions:
diagnosis
B    62.7 %
M    37.3 %
Name: proportion, dtype: str
```

```
In [5]: # Summary statistics for the 'mean' feature group
mean_cols = [c for c in df.columns if c.endswith('_mean')]
df[mean_cols].describe().round(3)
```

```
Out[5]:
```

	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean
count	569.000	569.000	569.000	569.000	569.000
mean	14.127	19.290	91.969	654.889	0.096
std	3.524	4.301	24.299	351.914	0.014
min	6.981	9.710	43.790	143.500	0.053
25%	11.700	16.170	75.170	420.300	0.086
50%	13.370	18.840	86.240	551.100	0.096
75%	15.780	21.800	104.100	782.700	0.105
max	28.110	39.280	188.500	2501.000	0.163

3. Data Preprocessing

Steps applied:

1. **Missing-value check** — verify the dataset before cleaning.
2. **Drop non-informative columns** — `id` (identifier) and `Unnamed: 32` (an all-NaN artifact column produced by a trailing comma in the CSV).
3. **Encode the target** — `M → 1`, `B → 0`.
4. **Stratified train/test split** (80/20) to preserve class proportions.
5. **Standardization** with `StandardScaler` — mandatory for SVM and Logistic Regression, harmless for Random Forest. To avoid **data leakage**, the scaler is fit on

the *training* data only (handled inside `Pipeline` objects during cross-validation, and explicitly for the held-out evaluation).

```
In [6]: # 1. Missing-value audit
missing = df.isnull().sum()
print('Columns with missing values:')
print(missing[missing > 0])
print('\n"Unnamed: 32" is entirely empty:', bool(df['Unnamed: 32'].isnull().sum() == 0))
print('Missing values in the 30 real features + target:',
      int(df.drop(columns=['id', 'Unnamed: 32']).isnull().sum().sum()))
```

Columns with missing values:

Unnamed: 32 569

dtype: int64

"Unnamed: 32" is entirely empty: True

Missing values in the 30 real features + target: 0

```
In [7]: # 2-3. Drop non-informative columns and encode the target
df = df.drop(columns=['id', 'Unnamed: 32'])
df['diagnosis'] = df['diagnosis'].map({'M': 1, 'B': 0})

print('Clean shape:', df.shape)
print('Encoded target counts:\n', df['diagnosis'].value_counts().rename({0: 'Benign', 1: 'Malignant'})
```

Clean shape: (569, 31)

Encoded target counts:



diagnosis

Benign(0) 357

Malignant(1) 212

Name: count, dtype: int64


```
In [8]: # Visualize target distribution
fig, axes = plt.subplots(1, 2, figsize=(12, 5))
df['diagnosis'].value_counts().sort_index().plot(kind='bar', ax=axes[0], color='steelblue')
axes[0].set_title('Diagnosis Distribution')
axes[0].set_xticklabels(['Benign (0)', 'Malignant (1)'], rotation=0)
axes[0].set_ylabel('Count')
axes[1].pie(df['diagnosis'].value_counts().sort_index(), labels=['Benign', 'Malignant'],
            autopct='%1.1f%%', colors=['steelblue', 'salmon'])
axes[1].set_title('Diagnosis Ratio')
plt.tight_layout()
plt.savefig('outputs/diagnosis_distribution.png', dpi=150, bbox_inches='tight')
plt.show()
```



```
In [9]: # Feature distributions: Benign vs Malignant (mean features)
fig, axes = plt.subplots(2, 5, figsize=(20, 8))
axes = axes.flatten()
for i, feature in enumerate(mean_cols):
    df[df['diagnosis'] == 0][feature].hist(ax=axes[i], alpha=0.6, color='steelblue')
    df[df['diagnosis'] == 1][feature].hist(ax=axes[i], alpha=0.6, color='salmon')
    axes[i].set_title(feature)
    axes[i].legend()
plt.suptitle('Feature Distributions: Benign vs Malignant', fontsize=14, y=1.05)
plt.tight_layout()
plt.savefig('outputs/feature_distributions.png', dpi=150, bbox_inches='tight')
plt.show()
```



```
In [10]: # Correlation heatmap (lower triangle)
plt.figure(figsize=(18, 14))
correlation = df.corr()
mask = np.triu(np.ones_like(correlation, dtype=bool))
sns.heatmap(correlation, mask=mask, annot=False, cmap='coolwarm', linewidths=1)
plt.title('Feature Correlation Matrix', fontsize=14)
plt.tight_layout()
plt.savefig('outputs/correlation_heatmap.png', dpi=150, bbox_inches='tight')
plt.show()
```



```
In [11]: # Features most correlated with the diagnosis
target_corr = df.corr()['diagnosis'].drop('diagnosis').abs().sort_values(ascending=False)
print('Top 10 features by |correlation| with diagnosis:')
print(target_corr.head(10).round(3))
```

Top 10 features by |correlation| with diagnosis:

concave points_worst	0.794
perimeter_worst	0.783
concave points_mean	0.777
radius_worst	0.776
perimeter_mean	0.743
area_worst	0.734
radius_mean	0.730
area_mean	0.709
concavity_mean	0.696
concavity_worst	0.660

Name: diagnosis, dtype: float64

4. Model Development

We split the data once into a stratified **80% train / 20% test** partition. All model selection (cross-validation and grid search) happens **only on the training set**; the test set is touched once, at the very end, for an unbiased performance estimate.

```
In [12]: X = df.drop(columns=['diagnosis'])
y = df['diagnosis']

X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=RANDOM_STATE, stratify=y)

print('Train set:', X_train.shape, '| Test set:', X_test.shape)
print('Train class balance:', dict(y_train.value_counts().sort_index()))
print('Test class balance:', dict(y_test.value_counts().sort_index()))

# Scaled arrays for the final held-out evaluation (scaler fit on train only)
scaler = StandardScaler().fit(X_train)
X_train_scaled = scaler.transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

```
Train set: (455, 30) | Test set: (114, 30)
Train class balance: {0: np.int64(285), 1: np.int64(170)}
Test class balance: {0: np.int64(72), 1: np.int64(42)}
```

4.1 Baseline Comparison with Stratified 5-Fold Cross-Validation

Before tuning, we compare the three candidate algorithms with stratified 5-fold cross-validation on the **training set**. Each model is wrapped in a `Pipeline` so that standardization is re-fit *inside every fold* — this prevents information from the validation fold leaking into scaling.

```
In [13]: cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=RANDOM_STATE)

candidates = {
    'Logistic Regression': LogisticRegression(max_iter=5000, random_state=RA
    'Random Forest':      RandomForestClassifier(n_estimators=100, random_s
    'SVM (RBF)':          SVC(kernel='rbf', probability=True, random_state=
}

cv_rows = []
for name, clf in candidates.items():
    pipe = Pipeline([('scaler', StandardScaler()), ('model', clf)])
    acc = cross_val_score(pipe, X_train, y_train, cv=cv, scoring='accuracy')
    f1 = cross_val_score(pipe, X_train, y_train, cv=cv, scoring='f1')
    cv_rows.append({'Model': name,
                    'CV Accuracy': f'{acc.mean():.4f} ± {acc.std():.4f}',
                    'CV F1': f'{f1.mean():.4f} ± {f1.std():.4f}'})

cv_results = pd.DataFrame(cv_rows).set_index('Model')
```

```
print('Stratified 5-fold cross-validation on the training set:')
cv_results
```

Stratified 5-fold cross-validation on the training set:

Out[13]:

	CV Accuracy	CV F1
Model		
Logistic Regression	0.9736 ± 0.0149	0.9640 ± 0.0207
Random Forest	0.9648 ± 0.0146	0.9519 ± 0.0207
SVM (RBF)	0.9714 ± 0.0054	0.9610 ± 0.0082

4.2 Hyperparameter Tuning (Grid Search)

We tune the two non-linear models with `GridSearchCV` (5-fold, optimizing F1).

Logistic Regression is kept as a fixed baseline. Grids are intentionally compact to keep run time short while still exploring the most influential parameters.

```
In [14]: # Random Forest grid
rf_pipe = Pipeline([('scaler', StandardScaler()),
                    ('model', RandomForestClassifier(random_state=RANDOM_STATE))])
rf_grid = {
    'model__n_estimators': [100, 200, 300],
    'model__max_depth': [None, 5, 10],
    'model__min_samples_split': [2, 5],
}
rf_search = GridSearchCV(rf_pipe, rf_grid, cv=cv, scoring='f1', n_jobs=-1)
rf_search.fit(X_train, y_train)
print('Random Forest -> best F1: %.4f' % rf_search.best_score_)
print('Best params:', rf_search.best_params_)
```

Random Forest -> best F1: 0.9553

Best params: {'model__max_depth': 5, 'model__min_samples_split': 2, 'model__n_estimators': 100}

```
In [15]: # SVM grid
svm_pipe = Pipeline([('scaler', StandardScaler()),
                    ('model', SVC(kernel='rbf', probability=True, random_state=RANDOM_STATE))])
svm_grid = {
    'model__C': [0.1, 1, 10, 100],
    'model__gamma': ['scale', 0.01, 0.1],
}
svm_search = GridSearchCV(svm_pipe, svm_grid, cv=cv, scoring='f1', n_jobs=-1)
svm_search.fit(X_train, y_train)
print('SVM (RBF) -> best F1: %.4f' % svm_search.best_score_)
print('Best params:', svm_search.best_params_)
```

SVM (RBF) -> best F1: 0.9674

Best params: {'model__C': 10, 'model__gamma': 'scale'}

4.3 Train Final Models

We train the three final models on the full training set: the tuned Random Forest, the tuned SVM, and the Logistic Regression baseline. Predictions and class probabilities are produced on the held-out test set for evaluation in the next section.

```
In [16]: logreg = Pipeline([('scaler', StandardScaler()),
                           ('model', LogisticRegression(max_iter=5000, random_state=
logreg.fit(X_train, y_train)

rf_best = rf_search.best_estimator_
svm_best = svm_search.best_estimator_

final_models = {'Logistic Regression': logreg,
                'Random Forest': rf_best,
                'SVM (RBF)': svm_best}

preds, probs = {}, {}
for name, model in final_models.items():
    preds[name] = model.predict(X_test)
    probs[name] = model.predict_proba(X_test)[:, 1]

for name in final_models:
    print(f'=== {name} ===')
    print(classification_report(y_test, preds[name], target_names=['Benign',
```

```
=== Logistic Regression ===
              precision    recall  f1-score   support

   Benign         0.96         0.99         0.97         72
  Malignant         0.97         0.93         0.95         42

 accuracy                   0.96         114
 macro avg         0.97         0.96         0.96         114
 weighted avg         0.97         0.96         0.96         114
```

```
=== Random Forest ===
              precision    recall  f1-score   support

   Benign         0.96         1.00         0.98         72
  Malignant         1.00         0.93         0.96         42

 accuracy                   0.97         114
 macro avg         0.98         0.96         0.97         114
 weighted avg         0.97         0.97         0.97         114
```

```
=== SVM (RBF) ===
              precision    recall  f1-score   support

   Benign         0.96         1.00         0.98         72
  Malignant         1.00         0.93         0.96         42

 accuracy                   0.97         114
 macro avg         0.98         0.96         0.97         114
 weighted avg         0.97         0.97         0.97         114
```

5. Model Evaluation

We evaluate every model on the **untouched test set** (114 patients) using Accuracy, Precision, Recall, F1-score and ROC-AUC, plus confusion matrices and ROC curves. Recall on the malignant class is highlighted as the clinically critical metric.


```
In [17]: rows = []
for name in final_models:
    p = preds[name]
    rows.append({
        'Model': name,
        'Accuracy': accuracy_score(y_test, p),
        'Precision': precision_score(y_test, p),
        'Recall': recall_score(y_test, p),
        'F1-Score': f1_score(y_test, p),
        'ROC-AUC': auc(*roc_curve(y_test, probs[name]))[:2]),
    })
results = pd.DataFrame(rows).set_index('Model').round(4)
print('Held-out test-set performance:')
results
```

Held-out test-set performance:


```
Out[17]:
```

	Accuracy	Precision	Recall	F1-Score	ROC-AUC
Model					
Logistic Regression	0.9649	0.975	0.9286	0.9512	0.9960
Random Forest	0.9737	1.000	0.9286	0.9630	0.9950
SVM (RBF)	0.9737	1.000	0.9286	0.9630	0.9927

```
In [18]: # Confusion matrices
fig, axes = plt.subplots(1, 3, figsize=(18, 5))
for ax, name in zip(axes, final_models):
    cm = confusion_matrix(y_test, preds[name])
    sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=ax,
                xticklabels=['Benign', 'Malignant'], yticklabels=['Benign',
                                                                    'Malignant'],
                ax.set_title(f'{name}\nConfusion Matrix')
                ax.set_ylabel('Actual'); ax.set_xlabel('Predicted'))
plt.tight_layout()
plt.savefig('outputs/confusion_matrices.png', dpi=150, bbox_inches='tight')
plt.show()
```




```
In [19]: # ROC curves
fig, ax = plt.subplots(figsize=(8, 6))
colors = ['seagreen', 'steelblue', 'salmon']
for (name, color) in zip(final_models, colors):
    fpr, tpr, _ = roc_curve(y_test, probs[name])
    ax.plot(fpr, tpr, color=color, lw=2, label=f'{name} (AUC = {auc(fpr, tpr)})')
ax.plot([0, 1], [0, 1], 'k--', lw=1.5, label='Random Classifier')
ax.set_xlabel('False Positive Rate'); ax.set_ylabel('True Positive Rate')
ax.set_title('ROC Curve Comparison'); ax.legend(loc='lower right'); ax.grid()
plt.tight_layout()
plt.savefig('outputs/roc_curve.png', dpi=150, bbox_inches='tight')
plt.show()
```



```
In [20]: # Feature importance from the tuned Random Forest
rf_importances = pd.Series(rf_best.named_steps['model'].feature_importances_,
                           index=X.columns).sort_values(ascending=False)
top15 = rf_importances.head(15)
```

```
plt.figure(figsize=(10, 6))
top15.plot(kind='bar', color='steelblue')
plt.title('Top 15 Feature Importances (Tuned Random Forest)')
plt.ylabel('Importance'); plt.xlabel('Feature')
plt.xticks(rotation=45, ha='right')
plt.tight_layout()
plt.savefig('outputs/feature_importance.png', dpi=150, bbox_inches='tight')
plt.show()
print('Top 5 features:')
print(top15.head(5).round(4))
```



Top 5 features:

area_worst	0.1529
concave points_worst	0.1345
radius_worst	0.1004
perimeter_worst	0.0876
concave points_mean	0.0837

dtype: float64

6. Interpretation of Results

Model comparison. All three models perform strongly on this well-separated dataset.

On the held-out test set the **tuned Random Forest** and **tuned SVM** are tied as the best models (accuracy = 0.974, precision = 1.000, F1 = 0.963), while the **Logistic Regression** baseline is only slightly behind (accuracy = 0.965). Notably, **all three models share the same recall on the malignant class (0.929)** — each misses the same 3 of 42 malignant cases — and the simple Logistic Regression actually attains the **highest ROC-AUC (0.996)**. This small spread is itself an important finding: the diagnostic signal in these FNA features is so strong and nearly linearly separable that even a basic linear model competes with the more complex non-linear ones. We therefore recommend the **tuned**

Random Forest as the final model — it matches the SVM on every test metric while also providing interpretable feature importances.

Cross-validation vs. test set. The stratified 5-fold CV scores (Section 4.1, accuracy ≈ 0.96 – 0.97) closely match the held-out test scores, confirming the results are **not an artifact of a lucky split** and that the models generalize. The very low CV standard deviations (e.g. SVM accuracy 0.9714 ± 0.0054) further indicate stable, reliable performance across folds.

Which features matter. The Random Forest feature-importance ranking is dominated by the **"worst" (largest) concave-points, radius, perimeter and area** measurements, together with their "mean" counterparts (top feature: `concave_points_worst`). Clinically this is intuitive: malignant nuclei tend to be **larger, more irregular in outline, and have more/deeper concave contour points** than benign nuclei. This matches the correlation analysis in Section 2, where the very same measurements (`concave_points_worst` , `concave_points_mean` , `radius_worst`) showed the strongest association with the diagnosis — giving us confidence that the models learn a medically meaningful signal rather than noise.

Error analysis. The confusion matrices show the few remaining errors are almost entirely **false negatives** (malignant predicted benign) — the costliest error type in this domain. Because the models output well-calibrated probabilities, a deployed system could **lower the decision threshold** to trade a little precision for higher recall, or flag borderline-probability cases for a mandatory second human review.

Usefulness. A model of this quality could serve as a **decision-support tool** that prioritizes suspicious FNA samples for pathologist review — not a replacement for a clinician, but a fast, consistent second opinion that helps catch cases early.

Limitations & future work.

- The dataset is a single-source benchmark (569 patients); external validation on other populations and imaging centers is needed before any clinical claim.
- Many features are highly correlated (Section 2 heatmap); dimensionality reduction (e.g. PCA) or feature selection could yield a simpler, equally accurate model.
- Future work: probability-threshold optimization to maximize malignant recall, calibration analysis, and testing additional algorithms (gradient boosting, neural networks).

Conclusion. Using a rigorous, leakage-free pipeline with stratified cross-validation and hyperparameter tuning, we built models that classify breast tumors as malignant or benign with $\sim 97\%$ test accuracy and ~ 0.99 ROC-AUC. The tuned Random Forest is recommended as the final model for its top-tier accuracy combined with interpretable feature importances.